

---

# Address Space Layout Randomization (ASLR)

---

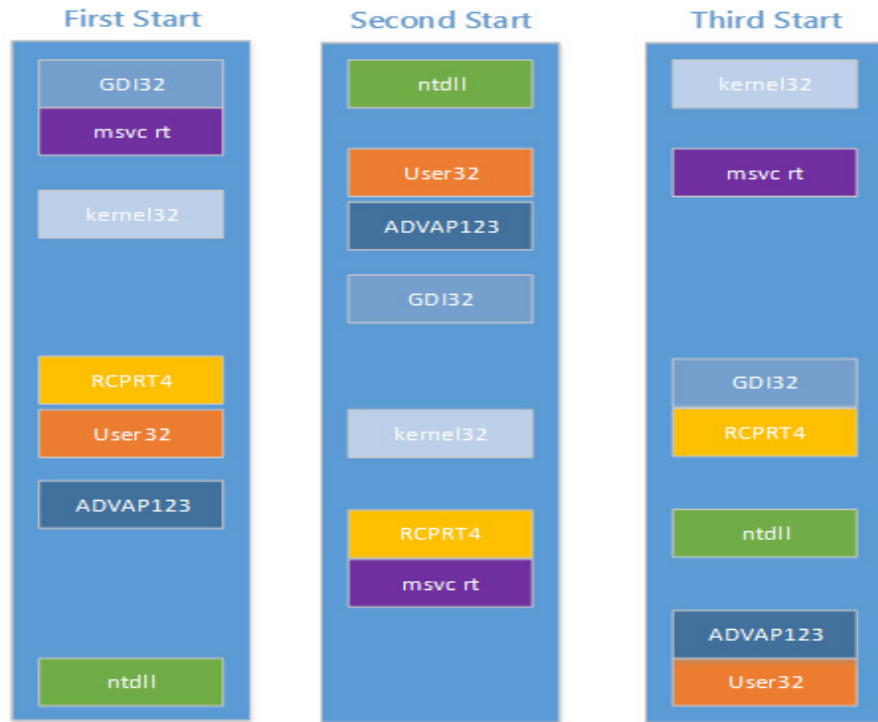
# The Attacker's Plan

- Find the bug in the source code (for eg. Kernel) that can be exploited
  - Eyeballing
  - Noticing something in the patches
  - Following CVE
- Use that bug to insert malicious code to perform something nefarious
  - Such as getting root privileges in the kernel

**Attacker depends upon knowing where these functions reside in memory. Assumes that many systems use the same address mapping. Therefore one exploit may spread easily.**

# Address Space Randomization

- Address space layout randomization (ASLR) randomizes the address space layout of the process
- Each execution would have a different memory map, thus making it difficult for the attacker to run exploits
- Initiated by Linux PaX project in 2001
- Now a default in many operating systems



Memory layout across boots for a Windows box

---

# ASLR in the Linux Kernel

- Locations of the base, libraries, heap, and stack can be randomized in a process' address space
- Built into the Linux kernel and controlled by `/proc/sys/kernel/randomize_va_space`
- `randomize_va_space` can take 3 values
  - 0** : disable ASLR
  - 1** : positions of stack, VDSO, shared memory regions are randomized  
the data segment is immediately after the executable code
  - 2** : (default setting) setting 1 as well as the data segment location is randomized

# ASLR in Action

```
chester@aahalya:~/tmp$ cat /proc/14621/maps
08048000-08049000 r-xp 00000000 00:15 81660111 /home/chester/tmp/a.out
08049000-0804a000 rw-p 00000000 00:15 81660111 /home/chester/tmp/a.out
b75da000-b75db000 rw-p 00000000 00:00 0
b75db000-b771b000 r-xp 00000000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771b000-b771c000 ---p 00140000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771c000-b771e000 r--p 00140000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771e000-b771f000 rw-p 00142000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771f000-b7722000 rw-p 00000000 00:00 0
b7734000-b7736000 rw-p 00000000 00:00 0
b7736000-b7737000 r-xp 00000000 00:00 0 [vdso]
b7737000-b7752000 r-xp 00000000 08:01 884950 /lib/ld-2.11.3.so
b7752000-b7753000 r--p 0001b000 08:01 884950 /lib/ld-2.11.3.so
b7753000-b7754000 rw-p 0001c000 08:01 884950 /lib/ld-2.11.3.so
bf9aa000-bf9bf000 rw-p 00000000 00:00 0 [stack]
```

First Run

```
chester@aahalya:~/tmp$ cat /proc/14639/maps
08048000-08049000 r-xp 00000000 00:15 81660111 /home/chester/tmp/a.out
08049000-0804a000 rw-p 00000000 00:15 81660111 /home/chester/tmp/a.out
b75dd000-b75de000 rw-p 00000000 00:00 0
b75de000-b771e000 r-xp 00000000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771e000-b771f000 ---p 00140000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b771f000-b7721000 r--p 00140000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b7721000-b7722000 rw-p 00142000 08:01 901176 /lib/i686/cmov/libc-2.11.3.so
b7722000-b7725000 rw-p 00000000 00:00 0
b7737000-b7739000 rw-p 00000000 00:00 0
b7739000-b773a000 r-xp 00000000 00:00 0 [vdso]
b773a000-b7755000 r-xp 00000000 08:01 884950 /lib/ld-2.11.3.so
b7755000-b7756000 r--p 0001b000 08:01 884950 /lib/ld-2.11.3.so
b7756000-b7757000 rw-p 0001c000 08:01 884950 /lib/ld-2.11.3.so
bfdd2000-bfde7000 rw-p 00000000 00:00 0 [stack]
```

Another Run

---

# ASLR in the Linux Kernel

- Permanent changes can be made by editing the `/etc/sysctl.conf` file

```
/etc/sysctl.conf, for example:  
kernel.randomize_va_space = value  
sysctl -p
```

---

# Internals : Making code relocatable

- **Load time relocatable**
  - where the loader modifies a program executable so that all addresses are adjusted properly
  - Relocatable code
    - Slow load time since executable code needs to be modified.
    - Requires a writeable code segment, which could pose problems
- **PIE : position independent executable**
  - a.k.a PIC (position independent code)
  - code that executes properly irrespective of its absolute address
  - Used extensively in shared libraries
    - Easy to find a location where to load them without overlapping with other modules

# Load Time Relocatable

1

```
unsigned long mylib_int;  
  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

```
chester@aahalya:~/sse/aslr$ make lib_reloc  
gcc -g -c mylib.c -o mylib.o  
gcc -shared -o libmylib.so mylib.o
```



Refer <https://chetrebeiro@bitbucket.org/casl/sse.git> (directory src/relocgot)

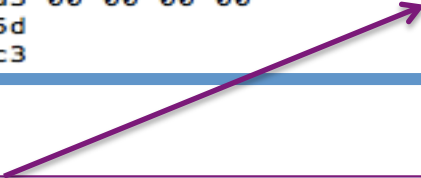


# Load Time Relocatable

```
unsigned long mylib_int;  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

2

```
0000046c <set_mylib_int>:  
46c: 55          push    %ebp  
46d: 89 e5       mov     %esp,%ebp  
46f: 8b 45 08    mov     0x8(%ebp),%eax  
472: a3 00 00 00 00 mov     %eax,0x0  
477: 5d         pop     %ebp  
478: c3         ret
```



note the 0x0 here...  
the actual address of mylib\_int is not filled in

# Load Time Relocatable

```
unsigned long mylib_int;  
  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

```
0000046c <set_mylib_int>:  
46c: 55          push    %ebp  
46d: 89 e5       mov     %esp,%ebp  
46f: 8b 45 08    mov     0x8(%ebp),%eax  
472: a3 00 00 00 mov     %eax,0x0  
477: 5d          pop     %ebp  
478: c3          ret
```

Relocatable table present in the executable that contains all references of mylib\_int

3

```
chester@aahalya:~/sse/aslr$ readelf -r libmylib.so  
  
Relocation section '.rel.dyn' at offset 0x304 contains 6 entries:  
Offset      Info      Type           Sym.Value    Sym. Name  
000015ec     00000008 R_386_RELATIVE 000015f8     mylib_int  
00000473     00000a01 R_386_32        000015f8     mylib_int  
0000047d     00000a01 R_386_32        000015f8     mylib_int  
000015cc     00000106 R_386_GLOB_DAT 00000000     __gmon_start__  
000015d0     00000206 R_386_GLOB_DAT 00000000     __Jv_RegisterClasses  
000015d4     00000306 R_386_GLOB_DAT 00000000     __cxa_finalize
```

Offset in memory where the fix needs to be made

Store binary value in the symbol memory location

CR

# Load Time Relocatable

```
unsigned long mylib_int;  
  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

```
0000046c <set_mylib_int>:  
46c: 55                push    %ebp  
46d: 89 e5             mov     %esp,%ebp  
46f: 8b 45 08          mov     0x8(%ebp),%eax  
472: a3 00 00 00 00    mov     %eax,0x0  
477: 5d                pop     %ebp
```

The loader fills in the actual address of mylib\_int at run time.

4 Breakpoint 1, main () at driver.c:9  
9 set\_mylib\_int(100);  
(gdb) disass set\_mylib\_int  
Dump of assembler code for function set\_mylib\_int:  
Relocation entries:  
Off 0xb7fde46c <set\_mylib\_int+0>: push %ebp  
Off 0xb7fde46d <set\_mylib\_int+1>: mov %esp,%ebp  
Off 0xb7fde46f <set\_mylib\_int+3>: mov 0x8(%ebp),%eax  
Off 0xb7fde472 <set\_mylib\_int+6>: mov %eax,0xb7fdf5f8  
Off 0xb7fde477 <set\_mylib\_int+11>: pop %ebp  
Off 0xb7fde478 <set\_mylib\_int+12>: ret  
End of assembler dump.  
Classes  
000015d4 00000306 R\_386\_GLOB\_DAT 00000000 \_\_cxa\_finalize

---

# Load Time Relocatable

## Limitations

- Slow load time since executable code needs to be modified
- Requires a writeable code segment, which could pose problems.
- Since executable code of each program needs to be customized, it would prevent sharing of code sections

---

# PIC Internals

- An additional level of indirection for all global data and function references
- Uses a lot of relative addressing schemes and a global offset table (GOT)
- For relative addressing,
  - data loads and stores should not be at absolute addresses but must be relative

Details about PIC and GOT taken from ...

<http://eli.thegreenplace.net/2011/11/03/position-independent-code-pic-in-shared-libraries/>

# Global Offset Table (GOT)

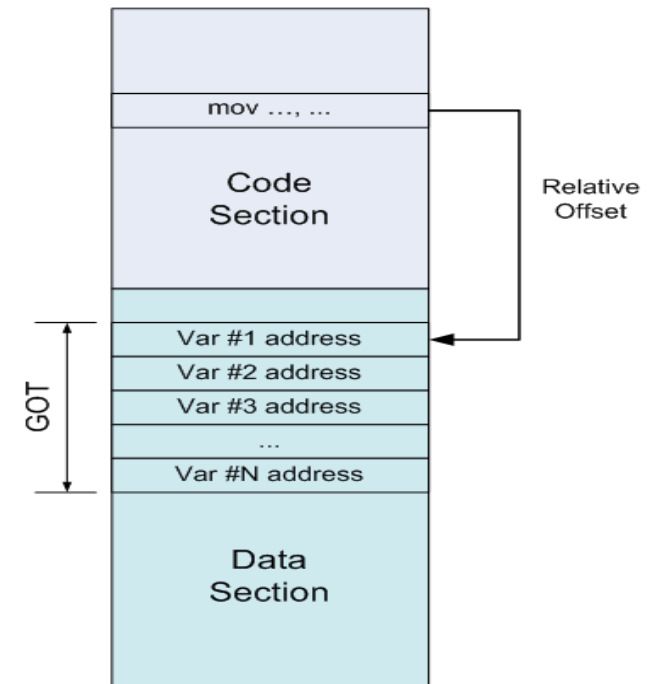
- Table at a fixed (known) location in memory space and known to the linker
- Has the location of the absolute address of variables and functions

## Without GOT

```
; Place the value of the variable in edx  
mov edx, [ADDR_OF_VAR]
```

## With GOT

```
; 1. Somehow get the address of the GOT into ebx  
lea ebx, ADDR_OF_GOT  
  
; 2. Suppose ADDR_OF_VAR is stored at offset 0x10  
;    in the GOT. Then this will place ADDR_OF_VAR  
;    into edx.  
mov edx, DWORD PTR [ebx + 0x10]  
  
; 3. Finally, access the variable and place its  
;    value into edx.  
mov edx, DWORD PTR [edx]
```



# Enforcing Relative Addressing (example)

```
unsigned long mylib_int;  
  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

With load time relocatable

```
0000046c <set_mylib_int>:  
46c: 55                push    %ebp  
46d: 89 e5            mov     %esp,%ebp  
46f: 8b 45 08        mov     0x8(%ebp),%eax  
472: a3 00 00 00 00  mov     %eax,0x0  
477: 5d              pop     %ebp  
478: c3              ret
```

With PIC

```
0000045c <set_mylib_int>:  
45c: 55                push    %ebp  
45d: 89 e5            mov     %esp,%ebp  
45f: e8 2b 00 00 00  call    48f <__i686.get_pc_thunk.cx>  
464: 81 c1 80 11 00 00 add     $0x1180,%ecx  
46a: 8b 81 f8 ff ff ff mov     -0x8(%ecx),%eax  
470: 8b 55 08        mov     0x8(%ebp),%edx  
473: 89 10          mov     %edx,(%eax)  
475: 5d              pop     %ebp  
476: c3              ret
```

```
0000048f <__i686.get_pc_thunk.cx>:  
48f: 8b 0c 24        mov     (%esp),%ecx  
492: c3              ret
```

# Enforcing Relative Addressing (example)

```
unsigned long mylib_int;  
  
void set_mylib_int(unsigned long x)  
{  
    mylib_int = x;  
}  
  
unsigned long get_mylib_int()  
{  
    return mylib_int;  
}
```

With load time relocatable

```
0000046c <set_mylib_int>:  
46c: 55                push    %ebp  
46d: 89 e5            mov     %esp,%ebp  
46f: 8b 45 08        mov     0x8(%ebp),%eax  
472: a3 00 00 00 00  mov     %eax,0x0  
477: 5d              pop     %ebp  
478: c3              ret
```

With PIC

Get address of next instruction  
to achieve relativeness

Index into GOT and get the  
actual address of mylib\_int into  
eax

Now work with the actual  
address.

```
0000045c <set_mylib_int>:  
45c: 55                push    %ebp  
45d: 89 e5            mov     %esp,%ebp  
45f: e8 2b 00 00 00  call    48f <__i686.get_pc_thunk.cx>  
464: 81 c1 80 11 00 00  add     $0x1180,%ecx  
46a: 8b 81 f8 ff ff ff  mov     -0x8(%ecx),%eax  
470: 8b 55 08        mov     0x8(%ebp),%edx  
473: 89 10          mov     %edx,(%eax)  
475: 5d              pop     %ebp  
476: c3              ret
```

```
0000048f <__i686.get_pc_thunk.cx>:  
48f: 8b 0c 24        mov     (%esp),%ecx  
492: c3              ret
```



---

# Advantage of the GOT

- With load time relocatable code, every variable reference would need to be changed
  - Requires writeable code segments
  - Huge overheads during load time
  - Code pages cannot be shared
- With GOT, the GOT table needs to be constructed just once during the execution
  - GOT is in the data segment, which is writeable
  - Data pages are not shared anyway
  - Drawback : runtime overheads due to multiple loads

---

# An Example of working with GOT

```
int myglob = 32;

int main(int argc, char **argv)
{
    return myglob + 5;
}
```

\$gcc -m32 -shared -fpic -S got.c

Besides a.out, this compilation also generates got.s

The assembly code for the program

```

.file      "got.c"
.globl myglob
.data
.align 4
.type     myglob, @object
.size     myglob, 4

myglob:
.long     32
.text
.globl main
.type     main, @function

main:
    pushl   %ebp
    movl    %esp, %ebp
    call    __i686.get_pc_thunk.cx
    addl     $ _GLOBAL_OFFSET_TABLE_, %ecx
    movl     myglob@GOT(%ecx), %eax
    movl     (%eax), %eax
    addl     $5, %eax
    popl     %ebp
    ret
    .size    main, .-main
    .ident   "GCC: (Debian 4.4.5-8) 4.4.5"
    .section .text.__i686.get_pc_thunk.cx,"axG",@progbits,__i686.get_
pc_thunk.cx,comdat
.globl __i686.get_pc_thunk.cx
.hidden __i686.get_pc_thunk.cx
.type     __i686.get_pc_thunk.cx, @function
__i686.get_pc_thunk.cx:
    movl     (%esp), %ecx
    ret
    .section .note.GNU-stack,"",@progbits

```

Data section

Text section

The macro for the GOT is known by the linker.  
%ecx will now contain the offset to GOT

Load the absolute address of myglob from the  
GOT into %eax

Fills %ecx with the eip of the next  
instruction.

Why do we need this indirect way of doing this?  
In this case what will %ecx contain?

# More

```
chester@aahalya:~/tmp$ readelf -S a.out
There are 27 section headers, starting at offset 0x69c:
```

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	00000000	000000	000000	00		0	0	0
[ 1]	.note.gnu.build-id	NOTE	000000d4	0000d4	000024	00	A	0	0	4
[ 2]	.hash	HASH								
[ 3]	.gnu.hash	GNU_HASH								
[ 4]	.dynsym	DYNAMIC								
[ 5]	.dynstr	STRTAB								
[ 6]	.gnu.version	VERSION								
[ 7]	.gnu.version_r	VERSION								
[ 8]	.rel.dyn	RELOCATION								
[ 9]	.rel.plt	RELOCATION								
[10]	.init	PROGBITS								
[11]	.plt	PROGBITS								
[12]	.text	PROGBITS								
[13]	.fini	PROGBITS								
[14]	.eh_frame	PROGBITS								
[15]	.ctors	PROGBITS								
[16]	.dtors	PROGBITS								
[17]	.jcr	PROGBITS								
[18]	.dynamic	DYNAMIC								
[19]	.got	PROGBITS								
[20]	.got.plt	PROGBITS								

```
chester@aahalya:~/tmp$ readelf -r ./a.out
Relocation section '.rel.dyn' at offset 0x2d8 contains 5 entries:
```

Offset	Info	Type	Sym.Value	Sym. Name
000015a8	00000008	R_386_RELATIVE		
00001584	00000106	R_386_GLOB_DAT	00000000	__gmon_start__
00001588	00000206	R_386_GLOB_DAT	00000000	Jv RegisterClasses
0000158c	00000406	R_386_GLOB_DAT	000015ac	myglob
00001590	00000306	R_386_GLOB_DAT	00000000	__cxa_finalize

offset of myglob in GOT

GOT it!

# Deep Within the Kernel

## (randomizing the data section)

loading the executable

```
1 static int load_elf_binary(struct linux_binprm *bprm, struct pt_regs *regs)
2 {
3     struct file *interpreter = NULL; /* to shut gcc up */
4     unsigned long load_addr = 0, load_bias = 0;
5     ...
6     #ifdef arch_randomize_brk
7         if ((current->flags & PF_RANDOMIZE) && (randomize_va_space > 1))
8             current->mm->brk = current->mm->start_brk =
9                 arch_randomize_brk(current->mm);
10    #endif
11    ...
12    out_free_ph:
13        kfree(elf_phdata);
14        goto out;
15
```

Check if randomize\_va\_space  
is > 1 (it can be 1 or 2)

```
1 unsigned long arch_randomize_brk(struct mm_struct *mm)
2 {
3     unsigned long range_end = mm->brk + 0x02000000;
4     return randomize_range(mm->brk, range_end, 0) ? : mm->brk;
5 }
```

Compute the end of the data  
segment (m->brk + 0x20)

```
10 unsigned long
11 randomize_range(unsigned long start, unsigned long end, unsigned long len)
12 {
13     unsigned long range = end - len - start;
14     if (end <= start + len)
15         return 0;
16     return PAGE_ALIGN(get_random_int() % range + start);
17 }
18
```

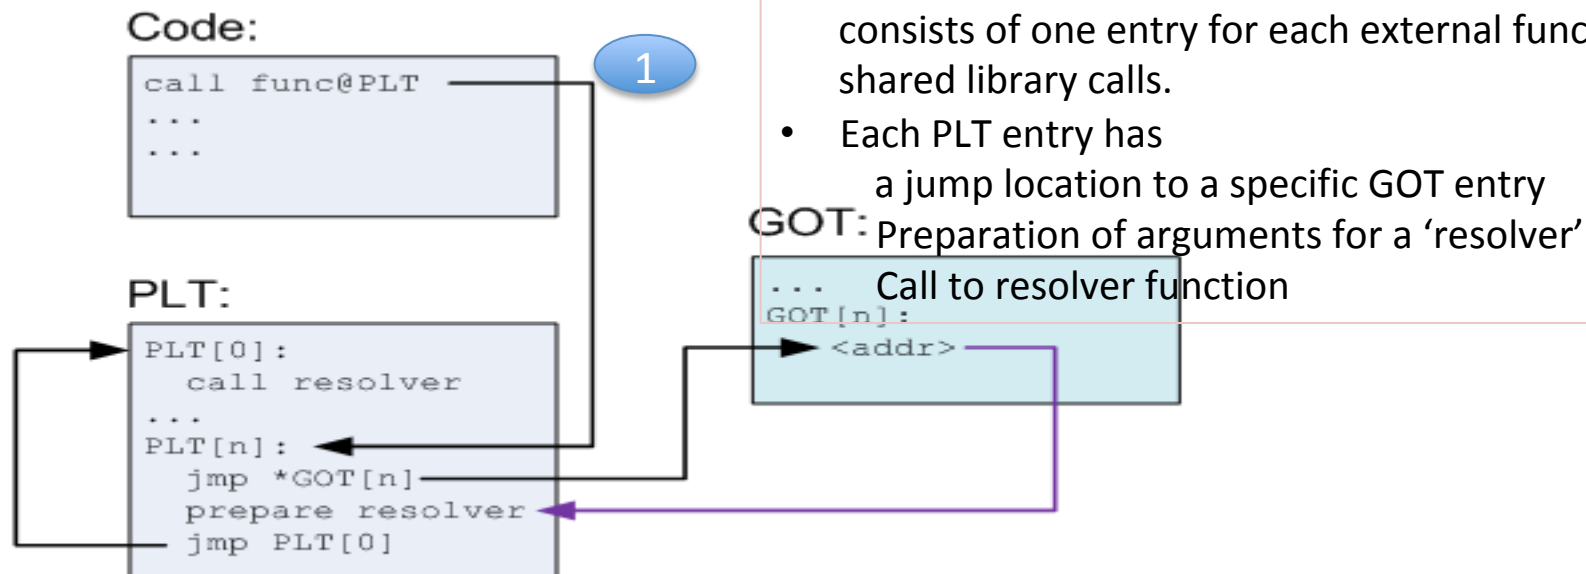
Finally Randomize

---

# Function Calls in PIC

- Theoretically could be done similar with the data...
  - call instruction gets location from GOT entry that is filled in during load time (this process is called binding)
  - In practice, this is time consuming. Much more functions than global variables. Most functions in libraries are unused
- Lazy binding scheme
  - Delay binding till invocation of the function
  - Uses a double indirection – PLT – **p**rocedure **l**inkage **t**able in addition to GOT

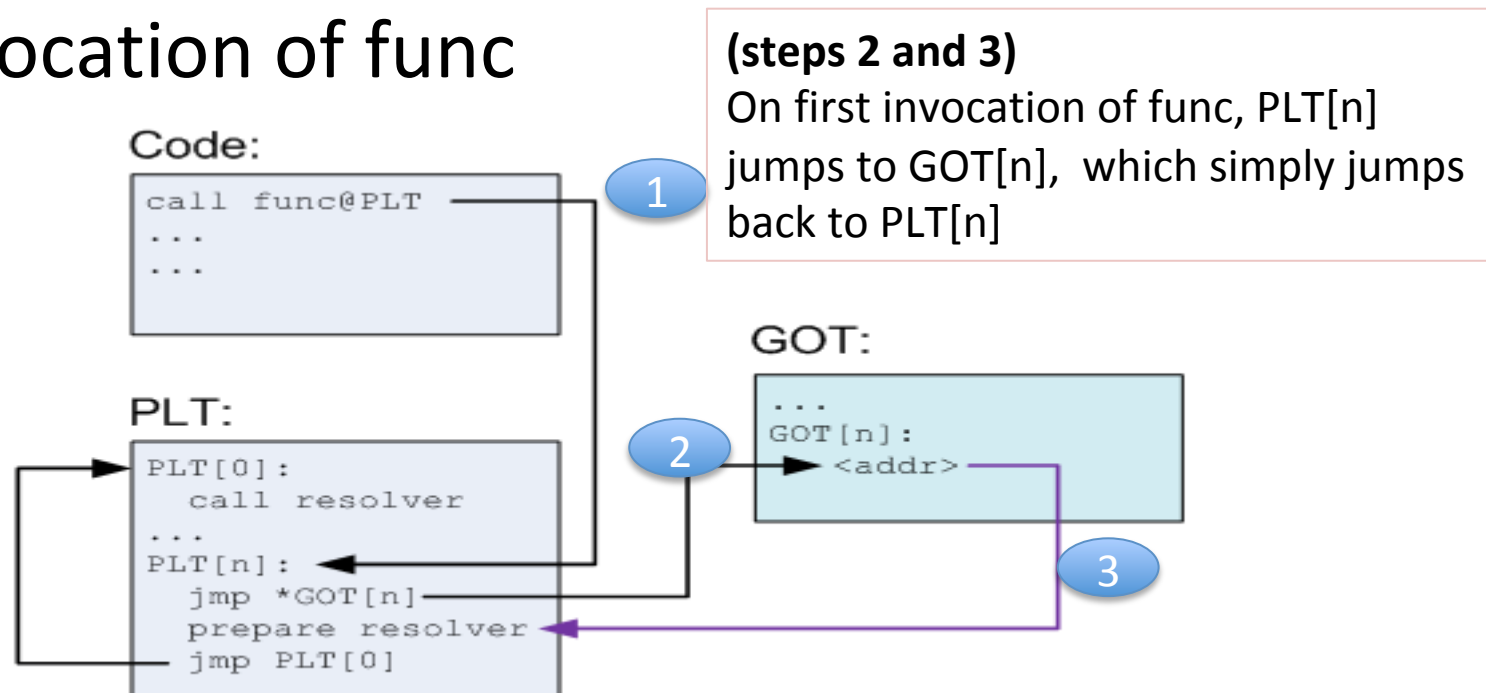
# The PLT



- Instead of directly calling func, invoke an offset in the PLT instead.
- PLT is part of the executable text section, and consists of one entry for each external function the shared library calls.
- Each PLT entry has
  - a jump location to a specific GOT entry

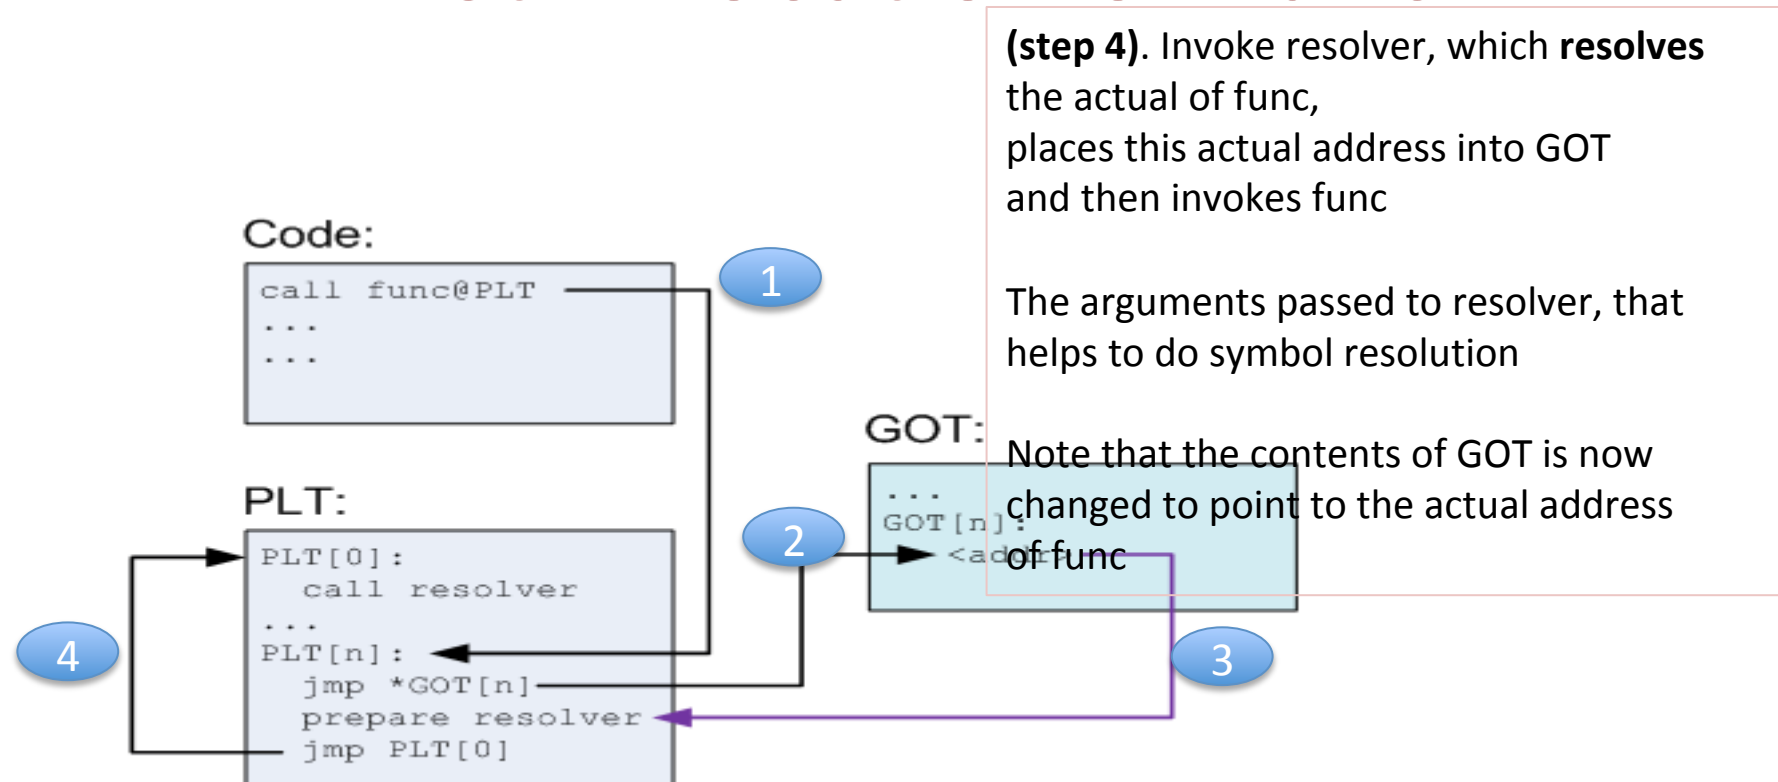
# First Invocation of Func

## First Invocation of func





# First Invocation of Func



# Example of PLT

```
unsigned long mylib_int;

void set_mylib_int(unsigned long x)
{
    mylib_int = x;
}

void inc_mylib_int()
{
    set_mylib_int(mylib_int + 1);
}

unsigned long get_mylib_int()
{
    return mylib_int;
}
```

```
chester@aahalya:~/sse/aslr/plt$ make
gcc -fpic -g -c mylib.c -o mylib.o
gcc -fpic -shared -o libmylib_pic.so mylib.o
```

Compiler converts the call to set\_mylib\_int into set\_mylib\_int@plt

```
000004b7 <inc_mylib_int>:
4b7: 55          push    %ebp
4b8: 89 e5       mov     %esp,%ebp
4ba: 53          push    %ebx
4bb: 83 ec 14    sub     $0x14,%esp
4be: e8 d4 ff ff ff call    497 <__i686.get_pc_thunk.bx>
4c3: 81 c3 81 11 00 00 add     $0x1181,%ebx
4c9: 8b 83 f8 ff ff ff mov     -0x8(%ebx),%eax
4cf: 8b 00       mov     (%eax),%eax
4d1: 83 c0 01    add     $0x1,%eax
4d4: 89 04 24    mov     %eax,(%esp)
4d7: e8 e0 fe ff ff call    3bc <set_mylib_int@plt>
4dc: 83 c4 14    add     $0x14,%esp
4df: 5b          pop     %ebx
4e0: 5d          pop     %ebp
4e1: c3          ret
```

# Example of PLT

Disassembly of section .plt:

```
0000039c <__gmon_start__@plt-0x10>:
39c:  ff b3 04 00 00 00      pushl  0x4(%ebx)
3a2:  ff a3 08 00 00 00      jmp     *0x8(%ebx)
3a8:  00 00                  add     %al, (%eax)
...

000003ac <__gmon_start__@plt>:
3ac:  ff a3 0c 00 00 00      jmp     *0xc(%ebx)
3b2:  68 00 00 00 00          push    $0x0
3b7:  e9 e0 ff ff ff         jmp     39c <_init+0x30>

000003bc <set_mylib_int@plt>:
3bc:  ff a3 10 00 00 00      jmp     *0x10(%ebx)
3c2:  68 08 00 00 00          push    $0x8
3c7:  e9 d0 ff ff ff         jmp     39c <_init+0x30>

000003cc <__cxa_finalize@plt>:
3cc:  ff a3 14 00 00 00      jmp     *0x14(%ebx)
3d2:  68 10 00 00 00          push    $0x10
3d7:  e9 c0 ff ff ff         jmp     39c <_init+0x30>
```

ebx points to the GOT table  
ebx + 0x10 is the offset  
corresponding  
to set\_mylib\_int

Offset of set\_mylib\_int in the  
GOT (+0x10).  
It contains the address of the  
next instruction (ie. 0x3c2)

```
chester@aahalya:~/sse/aslr/plt$ readelf -x .got.plt libmylib_pic.so
```

Hex dump of section '.got.plt':

```
0x00001644 6c150000 00000000 00000000 b2030000 l.....
0x00001654 c2030000 d2030000 .....
```

# Example of PLT

Disassembly of section .plt:

```
0000039c <__gmon_start__@plt-0x10>:
39c:  ff b3 04 00 00 00      pushl  0x4(%ebx)
3a2:  ff a3 08 00 00 00      jmp     *0x8(%ebx)
3a8:  00 00                  add     %al, (%eax)
...

000003ac <__gmon_start__@plt>:
3ac:  ff a3 0c 00 00 00      jmp     *0xc(%ebx)
3b2:  68 00 00 00 00 00      push    $0x0
3b7:  e9 e0 ff ff ff        jmp     39c <_init+0x30>

000003bc <set_mylib_int@plt>:
3bc:  ff a3 10 00 00 00      jmp     *0x10(%ebx)
3c2:  68 08 00 00 00 00      push    $0x8
3c7:  e9 d0 ff ff ff        jmp     39c <_init+0x30>

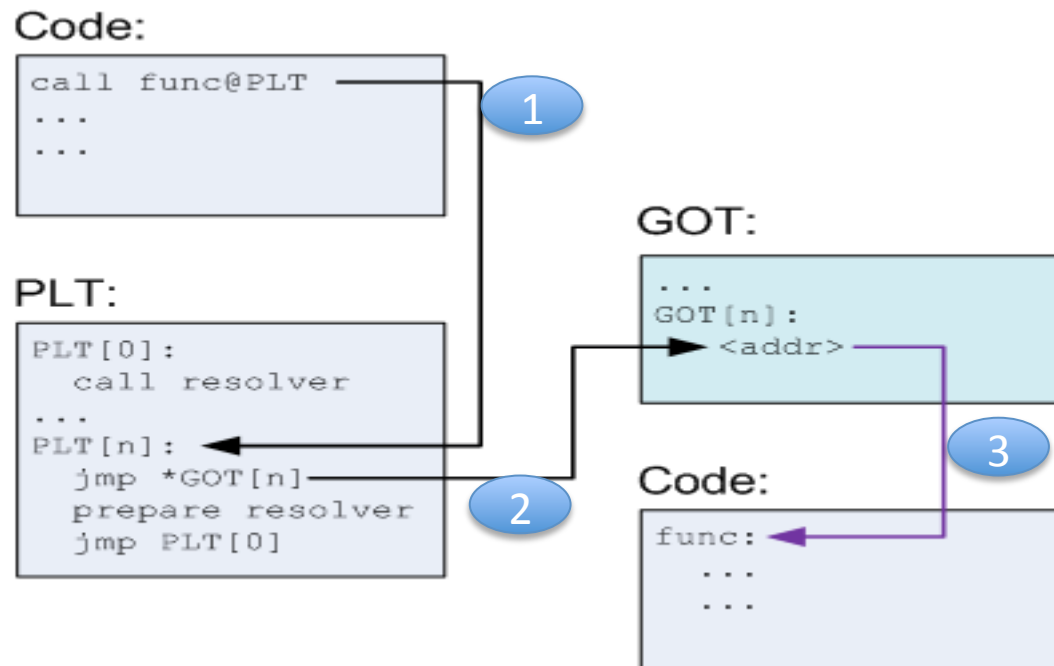
000003cc <__cxa_finalize@plt>:
3cc:  ff a3 14 00 00 00      jmp     *0x14(%ebx)
3d2:  68 10 00 00 00 00      push    $0x10
3d7:  e9 c0 ff ff ff        jmp     39c <_init+0x30>
```

Jump to the resolver, which resolves the actual address of set\_mylib\_int and fills it into the GOT

Push arguments for the resolver.

Jump to the first entry of the PLT ie. PLT[0]

# Subsequent invocations of Func



---

# Advantages

- Functions are relocatable, therefore good for ASLR
- Functions resolved only on need, therefore saves time during the load phase

---

# Bypassing ASLR

- Brute force
- Return-to-PLT
- Overwriting the GOT
- Timing Attacks